

CS2040 Tutorial 7 Suggested Solution

Week 9, starting 10 Oct 2022

Q1 Other BST Operations

You are given a BST implementation below:

```
class Node {
    int item;
    Node left, right;
    Node(int i, Node l, Node r) { item = i; left = l; right = r; }
}

class BST {
    int numNodes;
    Node root;

    int floor(int key) {} // to implement, may create helper method

    void insert(int key) { root = insert(key, root); }
    private Node insert(int key, Node curr) {
        if (curr == null) { numNodes++; return new Node(key, null, null); }
        if (key == curr.item) return curr; // no insertion
        if (key < curr.item) curr.left = insert(key, curr.left);
        else curr.right = insert(key, curr.right);
        return curr;
    }

    void preOrderPrint(Node root) {
        if (root == null) return;
        System.out.print(root.item + " ");
        preOrderPrint(root.left);
        preOrderPrint(root.right);
    }

    void inOrderPrint(Node root) {
        if (root == null) return;
        inOrderPrint(root.left);
        System.out.print(root.item + " ");
        inOrderPrint(root.right);
    }

    void postOrderPrint(Node root) {
        if (root == null) return;
        postOrderPrint(root.left);
        postOrderPrint(root.right);
        System.out.print(root.item + " ");
    }
}
```

```

void print() {
    System.out.print("Size: " + numNodes + "\nPreorder: [ ");
    preOrderPrint(root);
    System.out.print("]\nInorder: [ ");
    inOrderPrint(root);
    System.out.print("]\nPostorder: [ ");
    postOrderPrint(root);
    System.out.print("]\n");
}
}

```

(a) Write another method `int ceil(int key)` in the BST class that finds the *smallest* element that is more than or equals to `key`, or `Integer.MAX_VALUE` if none exists (i.e. the ceiling).

(b) Can `ceil(int)` be tweaked to implement `int higher(int key)` that returns the *smallest* element strictly greater than `key`, or `Integer.MAX_VALUE` if none exists? (i.e. the successor)

What is the time complexity of: one call to `higher(int)`, as well as repeated calls to `higher(currentKey)`, starting with `currentKey` being the smallest key in the tree, till we get `Integer.MAX_VALUE`?

(c) How does the BST implementation need to be changed to support `Node succ(Node curr)` (note the different parameter from (b)), that works similarly to `higher(int)` in (b) but returns the `Node` instead of the desired element?

Design and implement `succ(Node)`. Why is such an implementation more efficient?

Answer

(a) We can implement this recursively, similar to how binary search can be tweaked to allow similar `floor()`, or `bisect_left()` / `lower_bound()` (in py / C++ respectively) functionality.

When the target is $<$ a node's key, we visit the left subtree but still keep the current node's key as a candidate answer:

```

int ceil(int key) {
    Node ans = ceil(key, root);
    if (ans == null) return Integer.MAX_VALUE;
    return ans.item;
}

private Node ceil(int key, Node curr) {
    if (curr == null) return null;
    if (key == curr.item) return curr; // found ans, since no duplicates
    if (key > curr.item) return ceil(key, curr.right); // curr not valid
    // otherwise target key < curr.item, so curr is a candidate
    Node betterAns = ceil(key, curr.left);
    if (betterAns == null) return curr;
    return betterAns;
}

```

(b) Yes, the only difference is that when the target key is found in a Node's item, the Node is not a candidate solution:

```
int higher(int key) {
    Node ans = higher(key, root);
    if (ans == null) return Integer.MAX_VALUE;
    return ans.item;
}

private Node higher(int key, Node curr) {
    if (curr == null) return null;
    if (key >= curr.item) return higher(key, curr.right); //curr not valid
    // otherwise target key < curr.item, so curr is a candidate
    Node betterAns = higher(key, curr.left);
    if (betterAns == null) return curr;
    return betterAns;
}
```

One call to `higher(int)` runs in $O(h)$ time, h being the height of the BST.

When making repeated calls to `higher(int)` from smallest element to largest one key at a time, the algorithm moves down only 1 path from the root each call. Therefore it runs in a total of $O(Nh)$ time for a BST with N elements.

(c) Each Node needs to be augmented with a parent attribute. BST operations that change the structure of the tree need to maintain this attribute.

```
class BST {
    ...
    private Node insert(int key, Node curr) {
        if (curr == null) { numNodes++; return new Node(key, null, null); }
        if (key == curr.item) return curr; // no insertion
        if (key < curr.item) curr.left = insert(key, curr.left);
        else curr.right = insert(key, curr.right);
        if (curr.left != null) curr.left.parent = curr;
        if (curr.right != null) curr.right.parent = curr;
        return curr;
    }
}
```

Instead of starting from the root each time, we move from node to node, similar to how nodes are conceptually visited in an in-order traversal. When performing in-order traversal, the successor will be the smallest node in the right subtree if exists, otherwise it will be the first ancestor that turns right-ward.

```

class BST {
    ...
    Node succ(Node curr) { // pre-cond: curr not null
        if (curr.right != null) { // leftmost in right subtree
            Node smallest = curr.right;
            while (smallest.left != null) smallest = smallest.left;
            return smallest;
        }
        // no right subtree, so first right-ward ancestor
        Node ancestor = curr.parent;
        while (ancestor != null && ancestor.left != curr) {
            curr = ancestor;
            ancestor = ancestor.parent;
        }
        return ancestor;
    }
}

```

Therefore, while one call to `succ(Node)` still takes $O(h)$ time, N calls from the smallest key upward take a total of $O(N)$ time (better than $O(Nh)$ time previously).

Q2 Contiguous Strip with Same Colour

Given positive integers N and K , suppose you have created a computer game where there is a large $1 \times N$ strip of land painted black (colour = 0). The leftmost cell is indexed 0 while the rightmost cell is indexed $N-1$. Each cell has a colour value within 0 to K (a positive integer which could be $\gg N$) inclusive.

Design and implement a solution to perform the following operations, **each** running in $O(\log N)$ time or better:

```
void paint(int cell, int newColor)
```

paints the cell at the given index with a different new colour

```
int findContigLength(int cell)
```

returns the length of the longest contiguous sub-strip of land with the same colour that includes the (valid) cell with the given index

You may perform some initialization in $O(1)$ time.

e.g. $N = 5$, so the 5 cells start off with colours $[0, 0, 0, 0, 0]$. `findContigLength(1)` returns 5
 Next `paint(3, 5)` causes the colours to become $[0, 0, 0, 5, 0]$. `findContigLength(1)` returns 3
 Then `paint` is called 3 more times, giving colours $[0, 5, 5, 5, 5]$. `findContigLength(1)` returns 4
 Finally, `paint(3, 4)` is called, colours is now $[0, 5, 5, 4, 5]$. `findContigLength(1)` returns 2

Answer

Suppose we have the 5 cells shown in the above example, they can be represented as [(0, 0), (1, 0), (2, 0), (3, 0), (4, 0)] where each pair is a (index, colour) pair/object.

If we use an array to store this, paint() will be efficient but findContigLength() will be $O(A)$, where **A** is the answer of that operation. In the worst case, $O(A)$ can become $O(N)$.

We can just store strips that contain contiguous cells with the same colour instead, i.e. [(0, 4, 0)] on the initial state, and [(0, 0, 0), (1, 2, 5), (3, 3, 4), (4, 4, 5)] for the final state in the example. Each pair is now a (startIndex, endIndex, colour) triple/object. Once we find the desired strip, the strip length can be found in $O(1)$ time.

If the colours never change, we can use a sorted (by fromIndex) array and perform binary search. However, strips can be broken (added) and combined (removed) due to colour changes, so a sorted array becomes inefficient. Using a balanced BSTree set helps to achieve $O(\log N)$ insertions, removals and searches. However, there is a need to **maintain the invariant** that each strip of contiguous cells with the same colour should only appear in 1 element for the operation to be efficient.

To simplify the representation of the data, notice that each toIndex is always the next strip's fromIndex-1. Therefore, we just need to store (fromIndex, colour) pairs/objects [(0, 0), (1, 5), (3, 4), (4, 5)]. This saves coding time because we can use a balanced BSTree map instead of creating an object to represent the strip.

To prevent null values from appearing when querying near the first and last strip, we can insert dummy/sentinel strips before the first and last cells [(-1, ☹️), (0, 0), (1, 5), (3, 4), (4, 5), (5, ☹️)].

Operations:

Initialization creates an empty bBST and puts the 2 dummy strips into the bBST

findContigLength(cell) finds floor(cell) and higher(cell) and returns the distance between them

paint(cell, newColor) is tricky. It may create 1 or 2 new strips, and/or remove 1 or 2 existing strips. The trick is to find a way to make different decisions sequentially rather than consider all possibilities at once

```
class Strips {
    TreeMap<Integer, Integer> strips;

    Strips(int N, int K) { // K doesn't affect solution
        strips = new TreeMap<>(); strips.put(0, 0);
        strips.put(-1, -1); strips.put(N, -1); // dummy/sentinel strips
    }

    int findContigLength(int cell) {
        return strips.higherKey(cell) - strips.floorKey(cell);
    }
}
```

```

void paint(int cell, int newColor) {
    Integer prevColor = strips.lowerEntry(cell).getValue(),
        existingColor = strips.get(cell);

    strips.put(cell, newColor); // this cell, possibly replacing

    // 3 cases for breaking strips up:
    // (a) this cell was standalone, or at end of strip - do nothing
    // (b) not a standalone cell, and at middle of a strip
    if (1 + cell < strips.higherKey(cell) && existingColor == null)
        strips.put(1 + cell, prevColor); // next strip
    // (c) not a standalone cell, and at front of a strip
    else if (1 + cell < strips.higherKey(cell))
        strips.put(1 + cell, existingColor);

    // check if can combine prev strip
    if (prevColor == newColor)
        strips.remove(cell);

    // next (not mutually exclusive), check if can combine next strip
    Integer nextColor = strips.higherEntry(cell).getValue();
    if (newColor == nextColor)
        strips.remove(1 + cell);
}
}

```

Question 3 (**Online Discussion**) – Conditional Average

Modify the BST and or Node class in Q1, without changing the time complexity of any operation, such that it is possible to implement **double** findCondAverage(**int** upperBound) that returns, in $O(h)$ time, the average of all elements in the BST that are \leq upperBound.